

The slide features a decorative design with thin blue lines. A vertical line on the left and a horizontal line at the top intersect at the top-left corner, with a small blue quarter-circle ornament. Another horizontal line is positioned below the main text, and a vertical line on the right intersects it at the bottom-right corner, also with a small blue quarter-circle ornament.

# **Programming Shared Address Space Platforms (Pthreads)**

# Pthreads API

- ◆ The Pthreads API is defined in the ANSI/IEEE POSIX 1003.1 - 1995 standard.
- ◆ The subroutines which comprise the Pthreads API can be informally grouped into three major classes:
  - Thread management
  - Mutexes
  - Condition variables

# Thread Management

- ◆ The first class of functions work directly on threads - creating, detaching, joining, etc.
- ◆ They also include functions to set/query thread attributes (joinable, scheduling etc.)

# Creating Threads

- ◆ Initially, only a single, default thread. All other threads must be **explicitly created** by the programmer.
- ◆ **routine:** `pthread_create(thread, attr, start_routine, arg)`
  - **thread:** unique identifier for the new thread (*pthread\_t*)
  - **attr:** attribute object used to set thread attributes. (*pthread\_attr\_t*) You can specify a thread attributes object, or NULL for the default values.
  - **start\_routine:** C routine that the thread will execute.
  - **arg:** A single argument that may be passed to *start\_routine*. It must be passed by reference. NULL may be used if no argument is to be passed.
- ◆ If successful, the `pthread_create()` function shall return zero; otherwise, an error number shall be returned to indicate the error.

# Thread Attributes

- ◆ By default, a thread is created with certain attributes. Some of these attributes can be changed by the programmer via the thread attribute object.
- ◆ `pthread_attr_init(attr)` and `pthread_attr_destroy(attr)` are used to initialize/destroy the thread attribute object.
- ◆ Other routines are then used to query/set specific attributes in the thread attribute object.

# Terminating Thread

- ◆ There are several ways in which a Pthread may be terminated:
  - The thread makes a call to the `pthread_exit()` subroutine.
  - The thread is cancelled by another thread via the `pthread_cancel()` routine.
  - The entire process is terminated due to a call to the `exit` subroutine.

# Terminating Thread

- ◆ **routine:** `pthread_exit(status)`
- ◆ used to explicitly exit a thread.
- ◆ The programmer may optionally specify a termination status, which is stored as a void pointer for any thread that may join the calling thread.
- ◆ Cleanup: the `pthread_exit()` routine does not close files; any files opened inside the thread will remain open after the thread is terminated.

# Example

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid) {
    int *tid;
    tid = (int *) threadid;
    printf("Hello World! It's me, thread #%-d!\n", *tid);
    pthread_exit(NULL);
}
```



# Example

```
int main(int argc, char *argv[]) {  
    pthread_t threads[NUM_THREADS];  
    int rc, t, tids[NUM_THREADS];  
    for(t=0;t<NUM_THREADS;t++){  
        printf("In main: creating thread %d\n", t);  
        tids[t] = t;  
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &tids[t]);  
        if (rc){  
            printf("ERROR; return code from pthread_create() is %d\n", rc);  
            exit(-1);  
        }  
    }  
  
    pthread_exit(NULL);  
}
```

# Passing Arguments to Threads

- ◆ The `pthread_create()` routine permits the programmer to pass one argument to the thread start routine.
- ◆ For cases where multiple arguments must be passed, this limitation is easily overcome by creating a structure which contains all of the arguments, and then passing a pointer to that structure in the `pthread_create()` routine.
- ◆ All arguments must be passed by reference and cast to `(void *)`.

# Passing Arguments to Threads

```
struct two_args {  
    int arg1;  
    int arg2;  
};  
void *needs_2_args(void *ap) {  
    struct two_args *argp;  
    int a1, a2;  
  
    ...  
    argp = (struct two_args *) ap;  
    a1 = argp->arg1;  
    a2 = argp->arg2;  
  
    ...  
    free (argp);  
    pthread_exit(NULL);  
}
```

# Passing Arguments to Threads

```
int main(int argc, char *argv[]) {  
    pthread_t t;  
    struct two_args *ap;  
    int rc;  
  
    ...  
    ap = (struct two_args *) malloc (sizeof (struct two_args));  
    ap->arg1 = 1;  
    ap->arg2 = 2;  
    rc = pthread_create(&t, NULL, needs_2_args, (void *) ap);  
    ...  
    pthread_exit(NULL);  
}
```

# Joining & Detaching Threads

## ◆ Routines:

`pthread_join (threadid,status)`

`pthread_detach (threadid,status)`

`pthread_attr_setdetachstate (attr,detachstate)`

`pthread_attr_getdetachstate (attr,detachstate)`

- ◆ "Joining" is one way to accomplish synchronization between threads.
- ◆ The `pthread_join()` subroutine blocks the calling thread until the specified threadid thread terminates.
- ◆ The programmer is able to obtain the target thread's termination return status if it was specified in the target thread's call to `pthread_exit()`.

# Joining & Detaching Threads

- ◆ When a thread is created, one of its attributes defines whether it is joinable or detached.
- ◆ Only threads that are created as joinable can be joined. If a thread is created as detached, it can never be joined.
- ◆ To explicitly create a thread as joinable or detached, the attr argument in the `pthread_create()` routine is used:
  - Declare a pthread attribute variable of the `pthread_attr_t` data type
  - Initialize the attribute variable with `pthread_attr_init()`
  - Set the attribute detached status with `pthread_attr_setdetachstate()`
  - When done, free library resources used by the attribute with `pthread_attr_destroy()`

# Example

```
void *BusyWork(void *null) {  
    ...  
    pthread_exit((void *) 0);  
}
```

# Example

```
int main (int argc, char *argv[]) {
    pthread_attr_t attr;
    int rc, t;
    void *status;
    /* Initialize and set thread detached attribute */
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
    ...
    /* Free attribute and wait for the other threads */
    pthread_attr_destroy(&attr);
    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_join(thread[t], &status);
        ...
        printf("Completed join with thread %d status= %ld\n",t, (long)status);
    }
    pthread_exit(NULL);
}
```



# Synchronization Issues

- ◆ When multiple threads attempt to manipulate the same data item, the results can often be incoherent if proper care is not taken to synchronize them.
- ◆ Consider:

```
/* each thread tries to update variable best_cost as follows */  
if (my_cost < best_cost)  
    best_cost = my_cost;
```
- ◆ Assume that there are two threads, the initial value of `best_cost` is 100, and the values of `my_cost` are 50 and 75 at threads `t1` and `t2`.
- ◆ Depending on the schedule of the threads, the value of `best_cost` could be 50 or 75 - A **race condition** problem!

# Mutex

- ◆ The second class of functions deal with synchronization, called a "mutex", which is an abbreviation for "mutual exclusion".
- ◆ Mutex functions provide for creating, destroying, locking and unlocking mutexes.
- ◆ They are also supplemented by mutex attribute functions that set or modify attributes associated with mutexes.

# Creating & Destroying Mutex

## ◆ Routines:

pthread\_mutex\_init (*mutex*,*attr*)

pthread\_mutex\_destroy (*mutex*)

pthread\_mutexattr\_init (*attr*)

pthread\_mutexattr\_destroy (*attr*)

- ◆ Mutex must be declared with type pthread\_mutex\_t, and must be initialized before they can be used.

- ◆ There are two ways to initialize a mutex variable:

- Statically, when it is declared. For example:

```
pthread_mutex_t mymutex = PTHREAD_MUTEX_INITIALIZER;
```

- Dynamically, with the pthread\_mutex\_init() routine. This method permits setting mutex object attributes, *attr* (which may be specified as NULL to accept defaults).

- ◆ The mutex is initially unlocked.

# Locking & Unlocking Mutex

## ◆ Routines:

pthread\_mutex\_lock (*mutex*)

pthread\_mutex\_unlock (*mutex*)

pthread\_mutex\_trylock (*mutex*)

- ◆ pthread\_mutex\_lock() is used by a thread to acquire a lock on the specified *mutex* variable.
- ◆ pthread\_mutex\_unlock() will unlock a mutex if called by the owning thread. An error will be returned if:
  - the mutex was already unlocked
  - the mutex is owned by another thread
- ◆ pthread\_mutex\_trylock() will attempt to lock a mutex. However, if the mutex is already locked, the routine will return immediately with an "EBUSY" error code. This routine may be useful in preventing deadlock conditions.

# Example 1

◆ We can now write our previously incorrect code segment as:

```
pthread_mutex_t minimum_value_lock;
...
main() {
    ....
    pthread_mutex_init(&minimum_value_lock, NULL);
    ....
}
void *find_min(void *list_ptr) {
    ....
    pthread_mutex_lock(&minimum_value_lock);
    if (my_cost < best_cost)
        best_cost = my_cost;
    pthread_mutex_unlock(&minimum_value_lock);
}
```

# Example 2

- ◆ The **producer-consumer scenario** imposes the following constraints:
  - The producer thread must not overwrite the shared buffer when the previous task has not been picked up by a consumer thread.
  - The consumer threads must not pick up tasks until there is something present in the shared data structure.
  - Individual consumer threads should pick up tasks one at a time.

# Example 2

```
pthread_mutex_t task_queue_lock;  
int task_available;  
...  
main() {  
    ....  
    task_available = 0;  
    pthread_mutex_init(&task_queue_lock, NULL);  
    ....  
}
```

# Example 2

```
void *producer(void *producer_thread_data) {
    ....
    while (!done()) {
        inserted = 0;
        create_task(&my_task);
        while (inserted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (task_available == 0) {
                insert_into_queue(my_task);
                task_available = 1;
                inserted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
    }
}
```



# Example 2

```
void *consumer(void *consumer_thread_data) {  
    ...  
    while (!done()) {  
        extracted = 0;  
        while (extracted == 0) {  
            pthread_mutex_lock(&task_queue_lock);  
            if (task_available == 1) {  
                extract_from_queue(&my_task);  
                task_available = 0;  
                extracted = 1;  
            }  
            pthread_mutex_unlock(&task_queue_lock);  
        }  
        process_task(my_task);  
    }  
}
```

# Overheads of Locking

- ◆ Locks represent serialization points since critical sections must be executed by threads one after another.
- ◆ Encapsulating large segments of the program within locks can lead to significant performance degradation.
- ◆ It is often possible to reduce the idling overhead associated with locks using `pthread_mutex_trylock`.

# Alleviating Locking Overhead


```
/* using pthread_mutex_trylock routine */  
...  
Pthread_mutex_t tryLock_lock = PTHREAD_MUTEX_INITIALIZER;  
...  
lock_status=pthread_mutex_trylock(&tryLock_lock);  
if (lock_status == EBUSY) {  
    /* do something else */  
    ...  
}  
else {  
    /* do one thing */  
    ...  
    pthread_mutex_unlock(&tryLock_lock);  
}  
...
```

# Monitor

- ◆ Mutexs provide powerful synchronization tools. However,
  - Lock() and unlock() are scattered among several threads. Therefore, it is difficult to understand their effects
  - Usage must be correct in all the threads.
  - One bad thread (or one programming error) can kill the whole system.
- ◆ **Monitor** is a high-level abstraction that may provide a convenient and effective mechanism for thread synchronization


# Monitor

- ◆ Local data variables are accessible only by the monitor
- ◆ thread enters monitor by invoking one of its procedures
- ◆ Only one thread may be executing in the monitor at a time



# Monitor with Condition Variables

- ◆ Monitor does not need to code certain synchronization constraints explicitly.
- ◆ However, it is not sufficiently powerful for modeling some other synchronization schemes.
- ◆ An additional synchronization mechanism, i.e., **condition variable**, is required.



# Condition Variables

- ◆ The third class of functions address communications between threads that share a mutex.
- ◆ A condition variable allows a thread to block itself until specified data reaches a predefined state.
- ◆ A condition variable indicates an event and has no value.
  - One cannot store a value into nor retrieve a value from a condition variable.
  - If a thread must wait for an event to occur, that thread **waits** on the corresponding condition variable.
  - A condition variable has a queue for those threads that are waiting the corresponding event to occur to wait on.
  - If another thread causes the event to occur, that thread simply **signals** the corresponding condition variable.

# Condition Variables

- ◆ This class includes functions to create, destroy, wait and signal based upon specified variable values.
- ◆ Functions to set/query condition variable attributes are also included.
- ◆ A condition variable is **always used** in conjunction **with a mutex lock**.



# Creating & Destroying Condition Variables

## ◆ Routines:

`pthread_cond_init (condition, attr)`

`pthread_cond_destroy (condition)`

`pthread_condattr_init (attr)`

`pthread_condattr_destroy (attr)`

◆ Condition variables must be declared with type `pthread_cond_t`, and must be initialized before they can be used.

◆ There are two ways to initialize a condition variable:

- Statically, when it is declared. For example:  
`pthread_cond_t myconvar = PTHREAD_COND_INITIALIZER;`
- Dynamically, with the `pthread_cond_init()` routine. This method permits setting condition variable object attributes, `attr` (which may be specified as `NULL` to accept defaults).

# Waiting and Signaling on Condition Variables

## ◆ Routines:

pthread\_cond\_wait (*condition*, *mutex*)

pthread\_cond\_signal (*condition*)

pthread\_cond\_broadcast (*condition*)

- ◆ pthread\_cond\_wait() blocks the calling thread until the specified *condition* is signalled.
- ◆ This routine should be called while *mutex* is locked, and it will automatically release the *mutex* while it waits.
- ◆ After signal is received and thread is awakened, *mutex* will be automatically locked for use by the thread.
- ◆ The programmer is then responsible for unlocking *mutex* when the thread is finished with it.

# Waiting and Signaling on Condition Variables

- ◆ `pthread_cond_signal()` is used to signal (or wake up) another thread which is waiting on the condition variable.
  - It should be called after *mutex* is locked, and
  - must unlock *mutex* in order for `pthread_cond_wait()` routine to complete.
- ◆ The `pthread_cond_broadcast()` routine unlocks all of the threads blocked on the condition variable.

# Waiting and Signaling on Condition Variables

- ◆ Proper locking and unlocking of the associated *mutex* variable is essential when using these routines. For example:
  - Failing to lock the *mutex* before calling `pthread_cond_wait()` may cause it NOT to block.
  - Failing to unlock the *mutex* after calling `pthread_cond_signal()` may not allow a matching `pthread_cond_wait()` routine to complete (it will remain blocked).

# Producer-Consumer Using Condition Variables

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;
/* other data structures here */
...
main() {
    /* declarations and initializations */
    task_available = 0;
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
    ...
}
```

# Producer-Consumer Using Condition Variables

```
void *producer(void *producer_thread_data) {  
    while (!done()) {  
        create_task();  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 1)  
            pthread_cond_wait(&cond_queue_empty,  
                             &task_queue_cond_lock);  
        insert_into_queue();  
        task_available = 1;  
        pthread_cond_signal(&cond_queue_full);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
    }  
}
```

# Producer-Consumer Using Condition Variables

```
void *consumer(void *consumer_thread_data) {  
    while (!done()) {  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 0)  
            pthread_cond_wait(&cond_queue_full,  
                             &task_queue_cond_lock);  
        my_task = extract_from_queue();  
        task_available = 0;  
        pthread_cond_signal(&cond_queue_empty);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
        process_task(my_task);  
    }  
}
```

# Reference



POSIX Threads Programming, Lawrence Livermore National, Laboratory, <https://computing.llnl.gov/tutorials/pthreads/>